

# Challenges for Scheduling Scientific Workflows on Cloud Functions

Joanna Kijak, Piotr Martyna, Maciej Pawlik, Bartosz Balis  
and **Maciej Malawski**

Department of Computer Science,  
AGH University of Science and Technology  
al. Mickiewicza 30, 30-059 Kraków, Poland

# Outline

- Motivation: scientific workflows in clouds
- Experiments with HyperFlow
- Scheduling challenges
- Experiments with SDBWS algorithm
- Results on AWS Lambda
- Conclusions



# DICE Team

- Investigation of methods for building complex scientific collaborative applications
- Elaboration of environments and tools for e-Science
- Integration of large-scale distributed computing infrastructures
- Knowledge-based approach to services, components, and their semantic composition



**Other 15  
faculties**

**Department of Computer Science AGH (1980)**







800 students, 70 employees

<http://www.ki.agh.edu.pl/uk/index.htm>

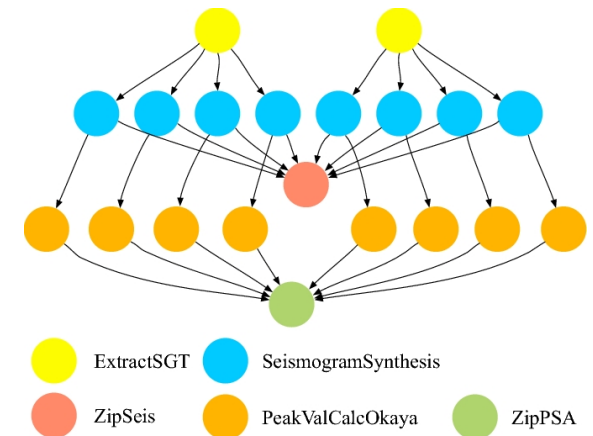


<http://dice.cyfronet.pl>

# Motivation: Scientific Workflows



- Astronomy, Geophysics, Genomics, Early Warning Systems ...
- Workflow = graph of tasks and dependencies, usually directed acyclic graph (DAG)
- Granularity of tasks
  - Large tasks (hours, days)
  - Small tasks (seconds, minutes)




# Infrastructure – from clusters to clouds

- Traditional HPC clusters in computing centers
  - Job scheduling systems
  - Local storage
- Grids to Clouds
  - Infrastructure as a service
  - Globally distributed
  - Virtual machines (VMs)
  - On-demand
  - Cost in \$\$ per time unit



# Workflow execution model in (traditional) clouds

- Workflow engine manages the tasks and dependencies
- Queue is used to dispatch ready tasks to the workers
- Worker nodes are deployed in Virtual Machines in the cloud
- Cloud storage such as Amazon S3 is used for data exchange
- Examples
  - **Pegasus**, Kepler, Triana, Pgrade, Askalon, ...
  - **HyperFlow** (AGH Krakow)





## Lightweight workflow programming and execution environment developed at AGH

### Programming language ecosystem



community

Text editor  
(favorite IDE, or  
generating script)

Wf graph




Wf activities

hflowc run  
workflow.json

### Workflow development

HyperFlow  
Engine

Executable workflow  
description (JSON)



### Workflow orchestration

Activity  $f_1$

Activity  $f_2$

Activity  $f_3$

VM



Local commands



Web  
Service

### Execution of workflow activities

App  
executables

HF Executor

Proxy  
cert.

### Simple wf description (JSON)

```
{
  "name": "PlotDataStatistics",
  "processes": [ {
    "name": "ComputeStats",
    "ins": [ "data.csv" ],
    "outs": [ "stats.txt" ],
    "config": {
      "executor": {
        "executable": "cstats.sh",
        "args": "data.csv -o stats.txt"
      }
    }
  }, {
    "name": "PlotChart",
    "ins": [ "stats.txt" ],
    "outs": [ "stats.png" ],
    "config": {
      "executor": {
        "executable": "plot.sh",
        "args": "stats.txt"
      }
    }
  }
  ]
}
```

### Advanced programming of wf activities (JavaScript)

```
function getPathWayByGene(ins, outs, config, cb) {
  var geneId = ins.geneId.data[0],
      url = ...

  http({ "timeout": 10000, "url": url },
    function(error, response, body) {
      ...
      cb(null, outs);
    });
}
```

Running a workflow – simple command line client


**hflowc run <workflow\_dir>**

**<workflow\_dir>** contains:

- File **workflow.json** (wf graph)
- File **workflow.cfg** (wf config)
- Optionally: file **functions.js** (advanced workflow activities)
- Input files

# New challenges – serverless architectures


- Serverless – no traditional VMs (servers)
- Composing of applications from existing cloud services
  - Typical example: web browser or mobile device interacting directly with the cloud
- Examples of services:
  - Databases: Firebase, DynamoDB
  - Messaging: Google Pub/Sub
  - Notification: Amazon SNS
- **Cloud Functions:**
  - Run a custom code on the cloud infrastructure





# Cloud Functions – good old RPC?

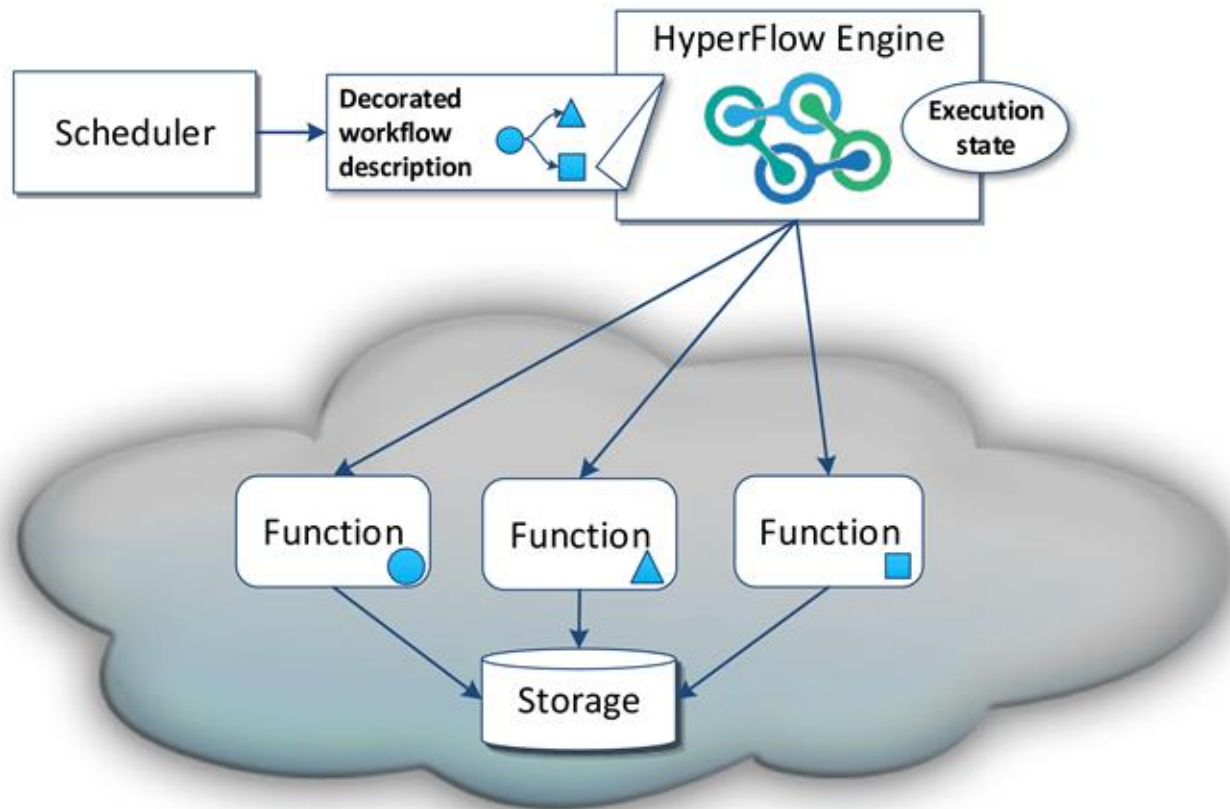
- Examples:
  - AWS Lambda
  - Google Cloud Functions (beta)
  - Azure Functions
  - IBM Bluemix OpenWhisk
- Functional programming approach:
  - Single function (operation)
  - **Not** a long-running service or process
  - Transient, stateless
- Infrastructure (execution environment) responsible for:
  - Startup
  - Parallel execution
  - Load balancing
  - Autoscaling
- Triggered by
  - Direct HTTP request
  - Change in cloud database
  - File upload
  - New item in the queue
  - Scheduled at specific time
- Developed in specific framework
  - Node.js, Java, Python
  - Custom code, libraries and binaries can be uploaded
- **Fine-grained pricing**
  - **Per 100ms \* GB (Lambda)**



# Scheduling challenges

- Which size of cloud functions should be allocated to each task of a workflow?
- Which tasks should be executed on FaaS and which ones on IaaS?
- What is the performance variability of the cloud functions infrastructure and how to deal with it?
- What are the limits of concurrency that we can expect when running multiple tasks as cloud functions in parallel?
- How to address the problem of data transfer between tasks?

# Execution model



[9] B. Baliś. Hyperflow: A model of computation, programming approach and enactment engine for complex distributed workflows

# Serverless Deadline-Budget Workflow Scheduling

- Adaptation of existing DBWS[2] heuristic for serverless model
- Finds mapping between tasks and resources (cloud functions) to meet the deadline constraint and tries to meet the budget constraint
- Assumes the knowledge of task runtime estimates on each resource type

# Idea of algorithm

- For each workflow level compute the maximum execution time

$$Level_{execution}^j = \max_{l(t_i)=j} \{ET_{max}(t_i)\}$$

- We divide the deadline into sub-deadlines for each level:

$$Level_{DL}^j = Level_{DL}^{j-1} + D_{user} * \frac{Level_{execution}^j}{\sum_{1 \leq j' \leq l(t_{exit})} Level_{execution}^{j'}}$$

# Resource selection

- Resource selection is based on the time and cost:

$$Time_Q(t_{cur}, r) = \frac{\xi * S_{DL}(t_{cur}) - FT(t_{cur}, r)}{FT_{max}(t_{cur}) - FT_{min}(t_{cur})}$$

$$Cost_Q(t_{cur}, r) = \frac{Cost_{max}(t_{cur}) - Cost(t_{cur}, r)}{Cost_{max}(t_{cur}) - Cost_{min}(t_{cur})} * \xi$$

$$\xi = \begin{cases} 1 & \text{if } FT(t_{cur}, r) < S_{DL}(t_{cur}) \\ 0 & \text{otherwise} \end{cases}$$



## Resource selection (2)

- We select the resource which maximizes the quantity:



$$Q(t_{cur}, r) = Time_Q(t_{cur}, r) * (1 - C_F) + Cost_Q(t_{cur}, r) * C_F$$

- Where  $C_F$  is a cost factor:

$$C_F = \frac{Cost_{low}(DAG)}{B_{user}}$$


- It represents user preferences:
  - Lower value means we prefer to pay more for faster execution
  - Higher value means we prefer cheaper and slower solutions

# Example DAG




# Task runtime estimates

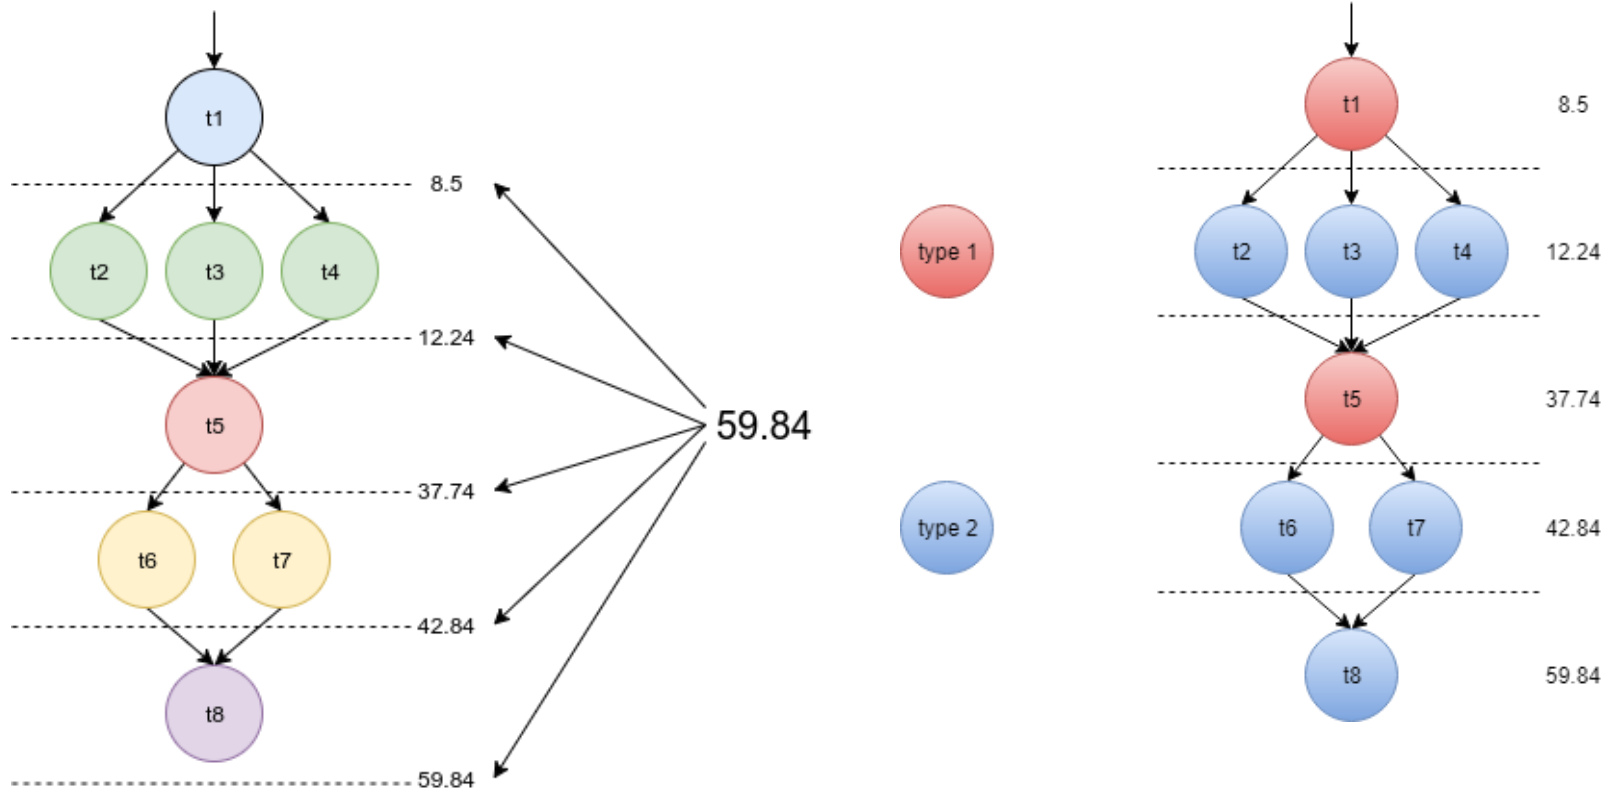
Function Type 1




Function Type 2




# Sub-deadlines and resource allocation



# Schedule



# Schedule







# Schedule



# Schedule






# Schedule



# Tests on AWS Lambda


- Montage workflow, 43 tasks
- Function size: 256, 512, 1024, 1536 MB
- Execution times estimated based on runs on homogeneous resources
- Limits adjusted to fit between minimum and maximum measured values

# Experiment 1





- Deadline: 18,6s
- Budget: \$0,00086
- 10,8% fastert han 1024MB
- 3,6% more expensive than 1024MB

# Experiment 2






- Deadline: 26,7s
- Budget: \$0,00094
- 30,9% faster than 512MB
- 7,4% more expensive than 512MB





# Experiment 3




- Deadline: 42,8s
- Budżet: \$0,00086
- 10,6% szybciej niż 512MB
- 4,9% drożej niż 512MB

# Conclusions

- Serverless and other highly-elastic infrastructures are interesting options for running high-throughput scientific workflows
- Cloud functions are heterogeneous
  - Technologies, APIs
  - Resource management policies (over/under provisioning)
  - Performance variations and guarantees
- Their provisioning model may change the game of resource management
- Experiments with SDBWS show that heterogeneous execution may have advantages, but more tests are needed

# Future Work

- Evaluation of parallelism limits and influence of delays
- Key parameter: elasticity
  - How quickly the infrastructure responds to the changes in workload demand
  - How fine-grained pricing can be?
  - Granularity of tasks vs. granularity of resources
- Example questions:
  - Which classes of tasks/workflows are suitable for such infrastructures?
  - How to dispatch tasks to various infrastructures?
  - Can we actually save costs when using such resources (e.g. for tight deadlines/high levels of parallelism)?



# Thank you!

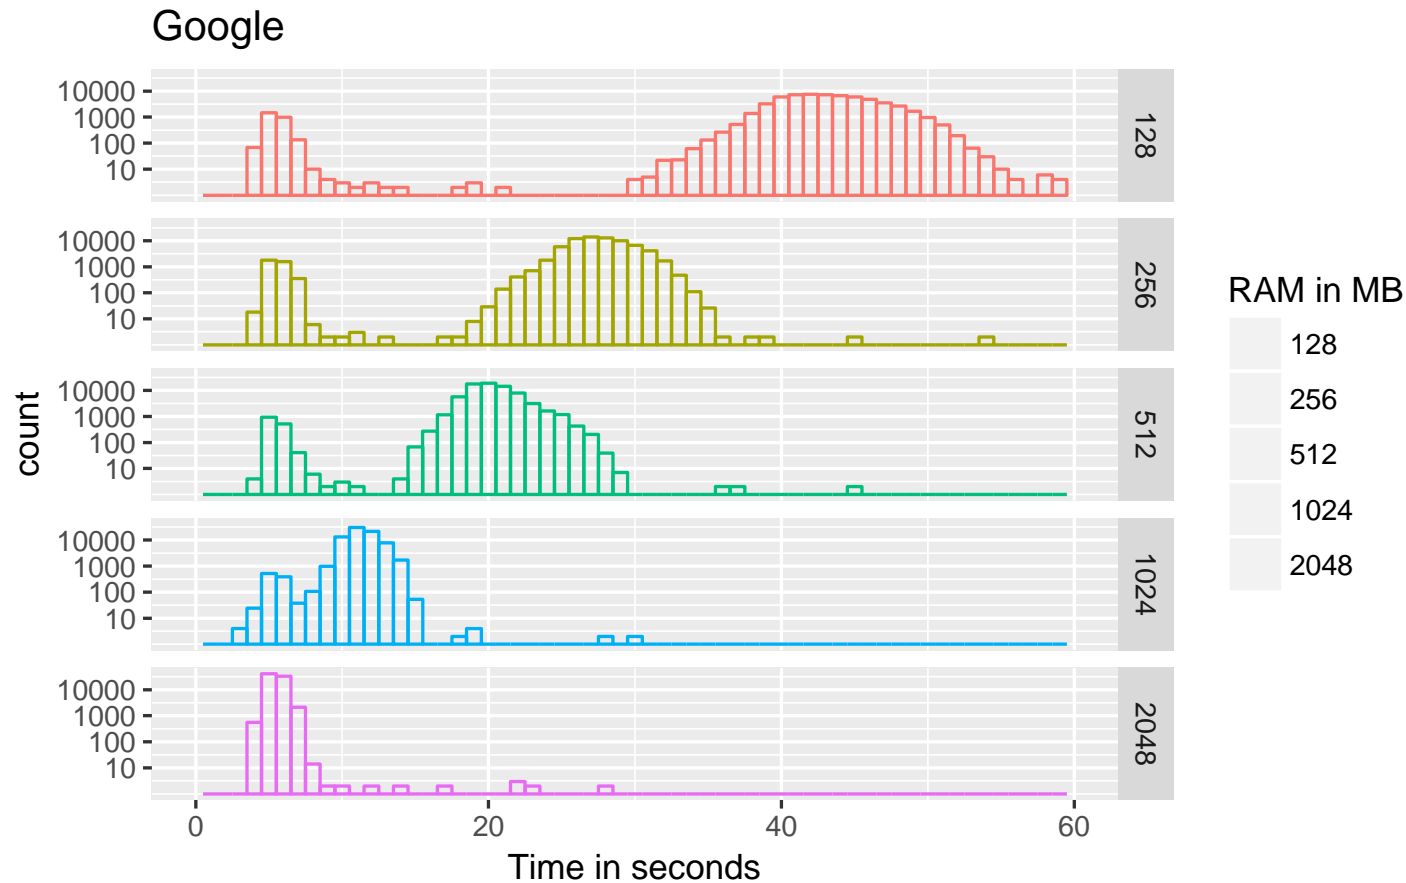
- DICE Team at AGH & Cyfronet
  - Marian Bubak, Piotr Nowakowski, Bartosz Baliś, Tomasz Gubała, Maciej Pawlik, Marek Kasztelnik, Bartosz Wilk, Tomasz Bartyński, Jan Meizner, Daniel Harężlak
- Collaboration
  - USC/ISI:
    - Ewa Deelman & Pegasus Team
  - Notre Dame:
    - Jarek Nabrzyski
- Projects & Grants
  - National Science Center (PL)
  - ISMOP (PL)
- References:
  - HyperFlow:  
<https://github.com/dice-cyfronet/hyperflow/>
  - DICE Team:  
<http://dice.cyfronet.pl>



# Backup slides

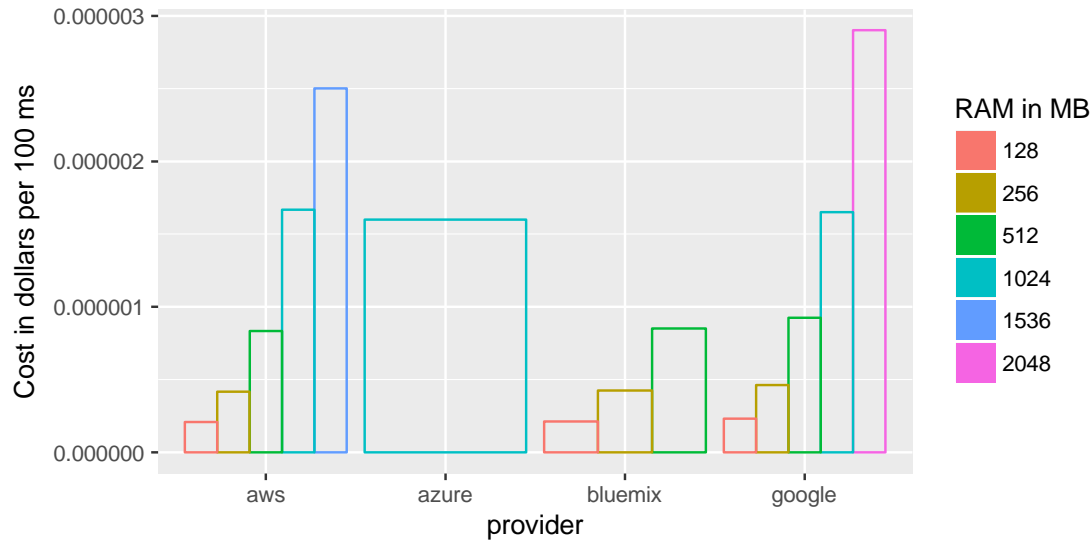
# Detailed Google Cloud Functions Performance Results

- Functions often run much faster than expected
- How often?  
About 5% times.






# Cost analysis



- List price vs. price/performance
- Different models:
  - AWS – proportional
  - IBM – invariant
  - Google: mixed
- For Azure we assume 1024 MB



# References

- [1] H. Arabnejad and J. G. Barbosa. List scheduling algorithm for heterogeneous systems by an optimistic cost table.
- [2] M. Ghasemzadeh, H. Arabnejad, and J. G. Barbosa. Deadline-budget constrained scheduling algorithm for scientific workflows in a cloud environment.
- [3] A. Ilyushkin, B. Ghit, and D. Epema. Scheduling workloads of workflows with unknown task runtimes.
- [4] M. Malawski, K. Figiela, M. Bubak, E. Deelman, and J. Nabrzyski. Scheduling multilevel deadline-constrained scientific workflows on clouds based on cost optimization.
- [5] M. Malawski, K. Figiela, A. Gajek, and A. Zima. Benchmarking heterogeneous cloud functions.
- [6] M. Malawski, K. Figiela, and J. Nabrzyski. Cost minimization for computational applications on hybrid cloud infrastructures.
- [7] M. Malawski, A. Gajek, A. Zima, and K. Figiela. Serverless execution of scientific workflows: Experiments with hyperflow, aws lambda and google cloud functions.
- [8] M. Malawski, G. Juve, E. Deelman, and J. Nabrzyski. Algorithms for cost- and deadline-constrained provisioning for scientific workflow ensembles in iaas clouds.
- [9] B. Baliś. Hyperflow: A model of computation, programming approach and enactment engine for complex distributed workflows