

# Comparison of FaaS Orchestration Systems



Pedro García López


**Cloud and Distributed Systems Lab**




UNIVERSITAT ROVIRA I VIRGILI

# CloudButton: Serverless Data Analytics

- 4.4M€ Research project
- cloudbutton.eu
- Coordinated by URV
- 2019-2021



# Creating Serverless Workflows




■ USER  
FUNCTIONS

■ WORKFLOW  
FRAMEWORK



AWS Step Functions



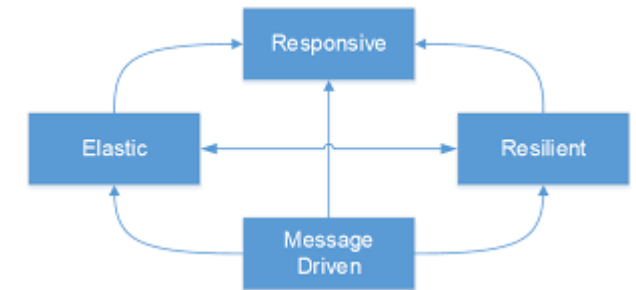
Azure Durable Functions



IBM Function Composer

# The Serverless Trilemma

- *If the serverless runtime is limited to a **reactive core**, i.e. one that deals only with dispatching functions in response to events, then these constraints form the serverless trilemma.*




- IBM Sequences are **ST-Safe**

- (1) Functions as Black Boxes
- (2) Substitution principle
- (3) Double billing

# Evaluation framework

- *ST-Safeness*
- *Programming model*
- *Parallel execution support*
- *State management*
- *Software packaging and repositories*
- *Architecture*
- *Billing model*
- *Overhead*



ComputerHope.com



AWS Step Functions

# Amazon Step Functions

➤ <i>ST-Safeness</i>	✗	(2) composability
➤ <i>Programming model</i>		Amazon States Language DSL
➤ <i>Parallel execution support</i>	✓	
➤ <i>State management</i>	32K	
➤ <i>Software packaging and repositories</i>		✓
➤ <i>Architecture</i>	client scheduler	
➤ <i>Billing model</i>	0.025USD per state transition	

# Amazon Step Functions



AWS Step Functions

```
StateMachine.Builder stateMachineBuilder =
    stateMachine()
        .comment("A state machine with par states.")
        .startAt("Parallel");

Branch.Builder[] branchBuilders =
    new Branch.Builder[NSTEPS];

for (int i = 0; i < NSTEPS; i++) {
    branchBuilders[i] = branch()
        .startAt(String.valueOf(i + 1))
        .state(String.valueOf(i + 1),
            taskState()
                .resource(arnTask).transition(end()));
}

stateMachineBuilder.state("Parallel",
    parallelState().branches(branchBuilders)
        .transition(end()));
final StateMachine stateMachine =
    stateMachineBuilder.build();
```



AWS Step Functions

# IBM Composer and Sequences

➤ <i>ST-Safeness</i>	✓	
➤ <i>Programming model</i>		JavaScript Composer library
➤ <i>Parallel execution support</i>	✗	
➤ <i>State management</i>		5MB
➤ <i>Software packaging and repositories</i>		✓
➤ <i>Architecture</i>	reactive core, conductor actions	
➤ <i>Billing model</i>	unknown, free?	



# IBM Composer and Sequences




```
composer.try(  
  composer.sequence(  
    'myWatsonTranslator/languageId',  
    composer.if(  
      p => p.language !== 'en',  
      composer.sequence(  
        p => ({translateFrom: p.language, translateTo: 'en', payload: p.payload}),  
        'myWatsonTranslator/translator'  
      ),  
      composer.sequence(  
        p => ({text: p.payload}),  
        'en2shakespeare'  
      )  
    )  
  ),  
  err => ({payload: 'Sorry we cannot translate your text'})  
)
```

# Azure Durable Functions

- *ST-Safeness* ✓
- *Programming model* C# async/await, Task Framework
- *Parallel execution support* ✓
- *State management* Unlimited, compressed
- *Software packaging and repositories* ✓
- *Architecture* reactive core, event sourcing
- *Billing model* unknown, unexpected storage costs

# Azure Durable Functions



```
public static async Task Run(DurableOrchestrationContext ctx)
{
    var parallelTasks = new List<Task<int>>();

    // get a list of N work items to process in parallel
    object[] workBatch = await ctx.CallActivityAsync<object[]>("F1");
    for (int i = 0; i < workBatch.Length; i++)
    {
        Task<int> task = ctx.CallActivityAsync<int>("F2", workBatch[i]);
        parallelTasks.Add(task);
    }

    await Task.WhenAll(parallelTasks);

    // aggregate all N outputs and send result to F3
    int sum = parallelTasks.Sum(t => t.Result);
    await ctx.CallActivityAsync("F3", sum);
}
```

# Experiment 1: Sequences




```
StateMachine.Builder stateMachineBuilder =  
    stateMachine()  
        .comment("A Sequence state machine")  
        .startAt("1");  
for (int i = 1; i <= NSTEPS; i++) {  
    stateMachineBuilder.state(String.valueOf(i),  
        taskState().resource(arnTask)  
        .transition((i != NSTEPS) ?  
            next(String.valueOf(i + 1)) : end()));  
}  
StateMachine stateMachine =  
    stateMachineBuilder.build();
```

```
composer.repeat(40, 'sleepAction')
```

```
for (int i = 0; i < NSTEPS; i++) {  
    await context.  
        CallActivityAsync("sleepAction", null);  
}
```



AWS Step Functions




# Experiment 2: Parallels

```
StateMachine.Builder stateMachineBuilder =
    stateMachine()
    .comment("A state machine with par. states.")
    .startAt("Parallel");

Branch.Builder[] branchBuilders =
    new Branch.Builder[NSTEPS];

for (int i = 0; i < NSTEPS; i++) {
    branchBuilders[i] = branch()
        .startAt(String.valueOf(i + 1))
        .state(String.valueOf(i + 1),
            taskState()
            .resource(arnTask).transition(end()));
}



stateMachineBuilder.state("Parallel",
    parallelState().branches(branchBuilders)
    .transition(end()));
final StateMachine stateMachine =
    stateMachineBuilder.build();
```



```
var tasks = new Task<long>[NSTEPS];
for (int i = 0; i < NSTEPS; i++)
{
    tasks[i] = context.CallActivityAsync<long>(
        "sleepAction");
}
await Task.WhenAll(tasks);
```

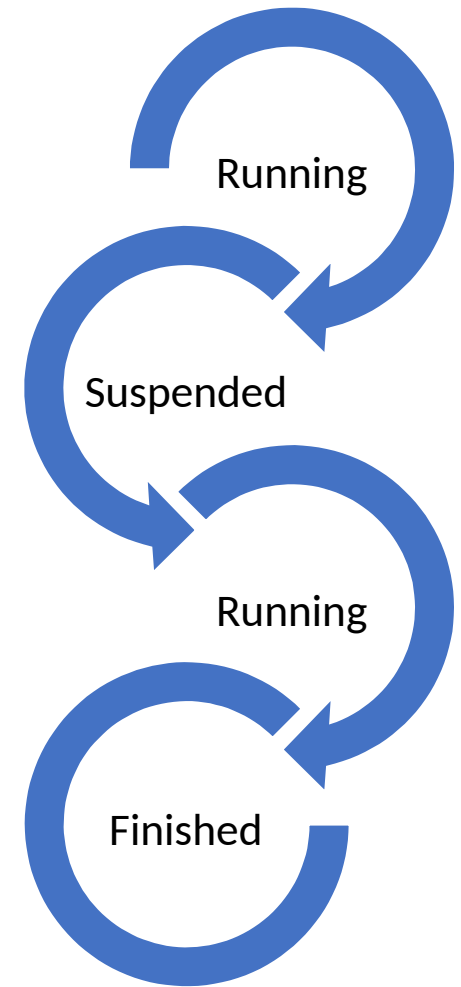


# Experiments



# Suspend API

- *Suspend function until event is received*
- *Passivation and state should be handled by the Function*
- *Requires a pure **reactive core** enabling **custom events***
- *It would enable the creation of custom orchestrators*



# Discussion

Metrics	Systems		
	<i>Amazon Step Functions</i>	<i>IBM Composer</i>	<i>Azure Durable Functions</i>
<i>ST-safe [1]</i>	<i>No</i> (compositions are not functions)	<i>Yes</i> (composition as functions)	<i>Yes</i> (composition as functions)
<i>Programming model</i>	DSL (JSON)	Composition library (Javascript)	async/await (C#)
<i>Reflective API</i>	<i>Yes (limited)</i>	<i>No</i>	<i>Yes</i>
<i>Parallel execution support</i>	<i>Yes (limited)</i>	<i>No</i>	<i>Yes (limited)</i>
<i>Software packaging and repositories</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes (no repo)</i>
<i>Billing model</i>	\$0.025 per 1,000 state transitions	Orchestrator function execution	Orchestrator function execution + storage costs
<i>Architecture</i>	Synchronous client scheduler	Reactive scheduler	Reactive scheduler





# Conclusions



- *Amazon Step Functions is the most mature project*
- *Microsoft ADF is the more advanced in programmability, IBM Composer wins in simplicity*
- *None of them support parallel tasks efficiently*
- *Orchestration must have a cost if it is fault-tolerant*
- *Reactive core, custom events and suspend API*
- *Early immature projects with high potential for the future*