

ACE: Just-in-time Serverless Software Component Discovery Through Approximate Concrete Execution

BY ANTHONY BYRNE¹, SHRIPAD NADGOWDA², AND AYSE K. COSKUN¹

¹Boston University; ²IBM T.J. Watson Research Center



Sixth International Workshop on Serverless Computing (WoSC6)
December 8, 2020

Serverless Containers: More Than Just FaaS

- “Serverless computing” encompasses more than Lambda functions
 - FaaS requirements (language, runtime, etc.) too strict for many developers
- Cloud providers offer serverless container platforms as a compromise
 - “Just hand us your Docker image, and we’ll handle everything else”
 - Bestow serverless benefits on any containerized app: scaling, billing, etc.
 - Allows executables not typically found in FaaS: compiled C/C++/Go binaries



AWS Fargate



Cloud Run

What Could Possibly Go Wrong?


- Cloud apps often made of in-house and off-the-shelf parts
 - Libraries, microservices, helper tools, etc.
- “Undesirable” software components key to many scandals
 - OpenSSL: Heartbleed bug exposed 66% of web servers (2014)
 - Apache Struts: 143 million Equifax records breached (2017)
 - Several court cases regarding licensing (e.g., AGPL)
- Binaries harder to screen for undesirables than Python/Java/JS code
 - No “requirements.txt” or other component manifest, just 1’s and 0’s



How can we discover software components in serverless binaries?

Just-in-time Component Discovery for Serverless

- Serverless gives cloud providers unprecedented access to developers' applications
 - Use it for good by scanning apps "JIT" before harm occurs
- Serverless binaries present special challenges
 - Metadata stripped and obscured through static linkage
 - Most analysis techniques slow and error-prone
- *Binary function fingerprinting* provides a framework
 - Disassemble binaries, fingerprint functions, check blocklist
 - If fingerprint similar to "known bad" one, then flag for review



How can we fingerprint a binary function?

Introducing Approximate Concrete Execution

```
push r1  
mov r2, 5  
jnz 0x56
```

Step 1

Disassemble raw binary and determine function bounds

- Prior work* provides function bounds in stripped binaries

```
sub r0, 8, r0  
stm r1, r0  
str 5, r2  
bisz r4, r3  
jcc r4, 56
```

Step 2

Translate assembly code to IR function-by-function

- REIL: simple MIPS-like register layout, infinite memory

REGISTERS	RAM
r0 = 10	0x0 = 0
r1 = 10	0x1 = 1
r2 = 10	0x2 = 2
r3 = 10	0x3 = 3
r4 = 10	0x4 = 4

Step 3

Filter code and provision a REIL “approximate VM”

- Remove all control flow instructions and sort (to account for compiler diffs.)

REGISTERS	RAM
r0 = 2	0x0 = 0
r1 = 10	0x1 = 1
r2 = 5	0x2 = 1
r3 = 10	0x3 = 3
r4 = 1	0x4 = 4

Step 4


Approximately execute and collect final aVM context

- Post-execution reg. values become fingerprint

* D. Andriesse, A. Slowinska, and H. Bos, “Compiler-Agnostic Function Detection in Binaries,” in *2017 IEEE European Symposium on Security and Privacy (EuroS&P)*, 2017, pp. 177–189.

Evaluating ACE for Serverless: Goals

- A JIT serverless binary fingerprinting method must...
 - produce representative fingerprints resistant to compiler variations
 - introduce very little overhead to the serverless environment
 - be tunable to different users' needs and applications



Evaluation


Accuracy: find the undesirable function

- Inject dummy function into ~230 C/C++ cloud apps
- Compile clean & injected apps and disassemble
- Classify each of the 37k functions using ACE
 - Positive: exact match to known dummy fingerprint


Overhead

- 5.2x faster end-to-end than baseline
 - Most functions fingerprinted in <10 milliseconds
- ACE requires no pre-training
 - Learning-based methods like SAFE require training and constant updating of ML models
- Minimal overall impact on cold-start or deployment latency

Experiment 1: Classification Accuracy



Experiment 2: Overhead Comparison



Why ACE for Serverless Component Discovery?




Speed

- No model training
- No complex instruction emulation
- 5.2x faster end-to-end



Resiliency

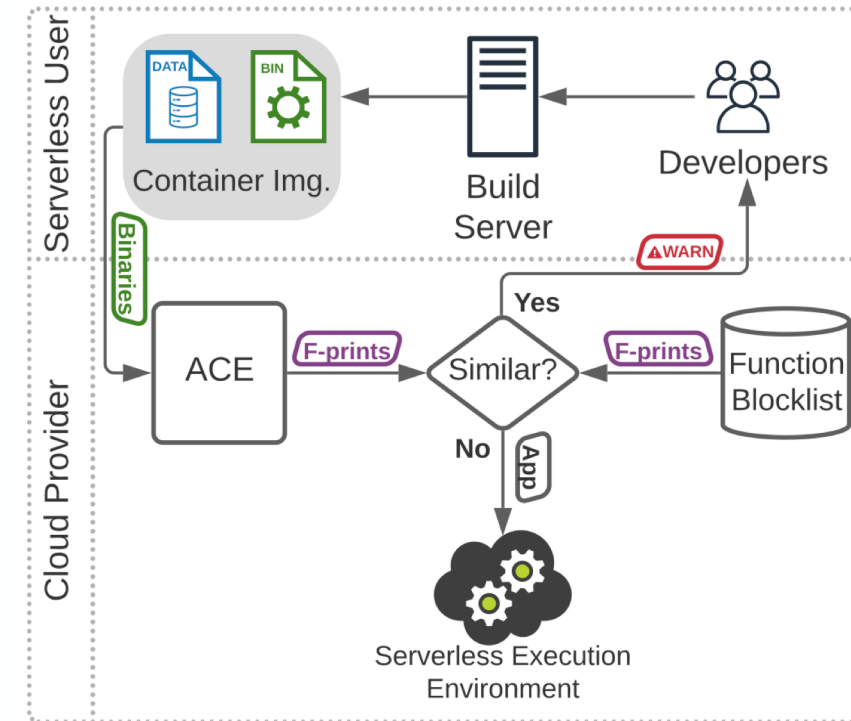
- Code filtering mitigates compiler variations
- 99% accurate binary classification of undesirable code



Versatility

- Several variables (sensitivity, aVM register size, etc.) tunable to users' needs
- Output vector suited to almost any search technique

- Serverless container platforms vulnerable to problems with **undesirable software components**
- ACE enables **just-in-time discovery** of these components through binary function fingerprinting
- We're excited to see future work apply the aVM to **further improve serverless performance & security**



Concluding Remarks

More info at bu.edu/peaclab

Please send feedback to abyrne19@bu.edu

This work has been partially funded by IBM Research